



Project: Finding the Best Markets to Advertise In

Introduction:

In this project I will be working with a fictitious e-learning company called Legends Code that offers courses on programming. Most of the courses offered are on web and mobile development, but they also cover other topics like data science, game development, etc. The e-learning company wants to promote their product and would like to invest money in advertisements.

The **goal** of this project is to find out the two best markets to advertise their product in. Given the cost of sending surveys to different target markets the company prefers to do some initial analysis using a free source.

In the project I will be using 2017 survey results from [FreeCodeCamp](#). FreeCodeCamp is a website that is funded by donations. Their main purpose is to help people learn to code for free. They accomplish this by creating videos, articles, and interactive coding lessons. Not only is FreeCodeCamp actually free but you can earn certificates as well. The survey data can be found here [GitHub Repository](#)

I will begin with loading in the libraries needed and initially explore the dataset.

```
In [3]: import pandas as pd
import numpy as np
import re
import matplotlib.pyplot as plt
plt.style.use('fivethirtyeight')
%matplotlib inline
# dataset has 136 rows. Avoiding the truncated view.
pd.options.display.max_columns=140

survey_2017 = pd.read_csv('2017-fCC-New-Coders-Survey-Data.csv', low_memory=False)
```

```
In [4]: survey_2017.head()
```

```
Out[4]:
```

	Age	AttendedBootcamp	BootcampFinish	BootcampLoanYesNo	BootcampName	BootcampRecommend	Childr
0	27.0	0.0	NaN	NaN	NaN	NaN	
1	34.0	0.0	NaN	NaN	NaN	NaN	

2	21.0	0.0	NaN	NaN	NaN	NaN
3	26.0	0.0	NaN	NaN	NaN	NaN
4	20.0	0.0	NaN	NaN	NaN	NaN

```
In [5]: survey_2017.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 18175 entries, 0 to 18174
Columns: 136 entries, Age to YouTubeTheNewBoston
dtypes: float64(105), object(31)
memory usage: 18.9+ MB
```

Exploring the columns:

Looking at the `survey.head` above to see what columns might be useful for my analysis. Age, CityPopulation, CountryLive, EmploymentField, EmploymentStatus, Gender, Income, JobRoleInterest, LanguageAtHome, MoneyForLearning and MonthsProgramming all stand out to me. They all seem like good preliminary columns to start exploring.

Since Legends Code's main teaching focus is web and app development I will start with the JobRoleInterest column.

Also a source of understanding each column can be found here. [FreeCodeCamp Github](#)

```
In [6]: print(survey_2017['JobRoleInterest'].isnull().value_counts())

round((survey_2017['JobRoleInterest'].isnull().value_counts() / 18175) * 100,2)

True      11183
False      6992
Name: JobRoleInterest, dtype: int64

Out[6]: True      61.53
False     38.47
Name: JobRoleInterest, dtype: float64
```

Comments:

As we can see above 61.53% of the column is a non null value. I have almost 39% that is missing. Lets see what the value counts are.

```
In [7]: survey_2017['JobRoleInterest'].value_counts().head(30)

Out[7]: Full-Stack Web Developer      823
        Front-End Web Developer      450
        Data Scientist               152
        Back-End Web Developer       142
        Mobile Developer             117
```

Game Developer	114
Information Security	92
Full-Stack Web Developer, Front-End Web Developer	64
Front-End Web Developer, Full-Stack Web Developer	56
Product Manager	55
Data Engineer	53
User Experience Designer	52
User Experience Designer, Front-End Web Developer	43
Front-End Web Developer, Back-End Web Developer, Full-Stack Web Developer	39
Back-End Web Developer, Full-Stack Web Developer, Front-End Web Developer	36
Back-End Web Developer, Front-End Web Developer, Full-Stack Web Developer	36
DevOps / SysAdmin	36
Full-Stack Web Developer, Front-End Web Developer, Back-End Web Developer	31
Front-End Web Developer, Full-Stack Web Developer, Back-End Web Developer	30
Full-Stack Web Developer, Mobile Developer	29
Front-End Web Developer, User Experience Designer	29
Back-End Web Developer, Full-Stack Web Developer	27
Full-Stack Web Developer, Back-End Web Developer	26
Back-End Web Developer, Front-End Web Developer	20
Data Engineer, Data Scientist	19
Full-Stack Web Developer, Back-End Web Developer, Front-End Web Developer	19
Front-End Web Developer, Mobile Developer	18
Full-Stack Web Developer, Data Scientist	17
Mobile Developer, Game Developer	16
Data Scientist, Data Engineer	16

Name: JobRoleInterest, dtype: int64

Comments:

Above I can see some survey respondents listed multiple jobs. For example we have Front-End Web Developer listed over 10 times spread across the entire top 30 value counts.

Below there is a column named JobInterestFrontEnd this column shows 4047 entries for the role Front End Web Developer.

I'll clean the Job Role Interest column first and see how many people are interested in multiple job rolls. Then I'll see how the job role interest compare to the job interest front end column

```
In [8]: survey_2017['JobInterestFrontEnd'].value_counts()
```

```
Out[8]: 1.0    4047
Name: JobInterestFrontEnd, dtype: int64
```

```
In [9]: survey_2017['JobRoleInterest'].astype(str)
```

```
survey_2017['JobRoleInterest'] = survey_2017['JobRoleInterest'].str.replace(r' |-', ' ')
```

```
In [10]: survey_2017['JobRoleInterest'].value_counts().head(30)
```

```
Out[10]: [full stack web developer]      823
[front end web developer]      450
[data scientist]               152
[back end web developer]      142
[mobile developer]            117
[game developer]              114
[information security]         92
[full stack web developer, front end web developer]  64
[front end web developer, full stack web developer]  56
[product manager]             55
[dataengineer]                53
[user experience designer]     52
```

[user experience designer, front end web developer]	43
[front end web developer, back end web developer, full stack web developer]	39
[back end web developer, full stack web developer, front end web developer]	36
[devops / sysadmin]	36
[back end web developer, front end web developer, full stack web developer]	36
[full stack web developer, front end web developer, back end web developer]	31
[front end web developer, full stack web developer, back end web developer]	30
[full stack web developer, mobile developer]	29
[front end web developer, user experience designer]	29
[back end web developer, full stack web developer]	27
[full stack web developer, back end web developer]	26
[back end web developer, front end web developer]	20
[data engineer, data scientist]	19
[full stack web developer, back end web developer, front end web developer]	19
[front end web developer, mobile developer]	18
[full stack web developer, data scientist]	17
[mobile developer, game developer]	16
[data scientist, data engineer]	16

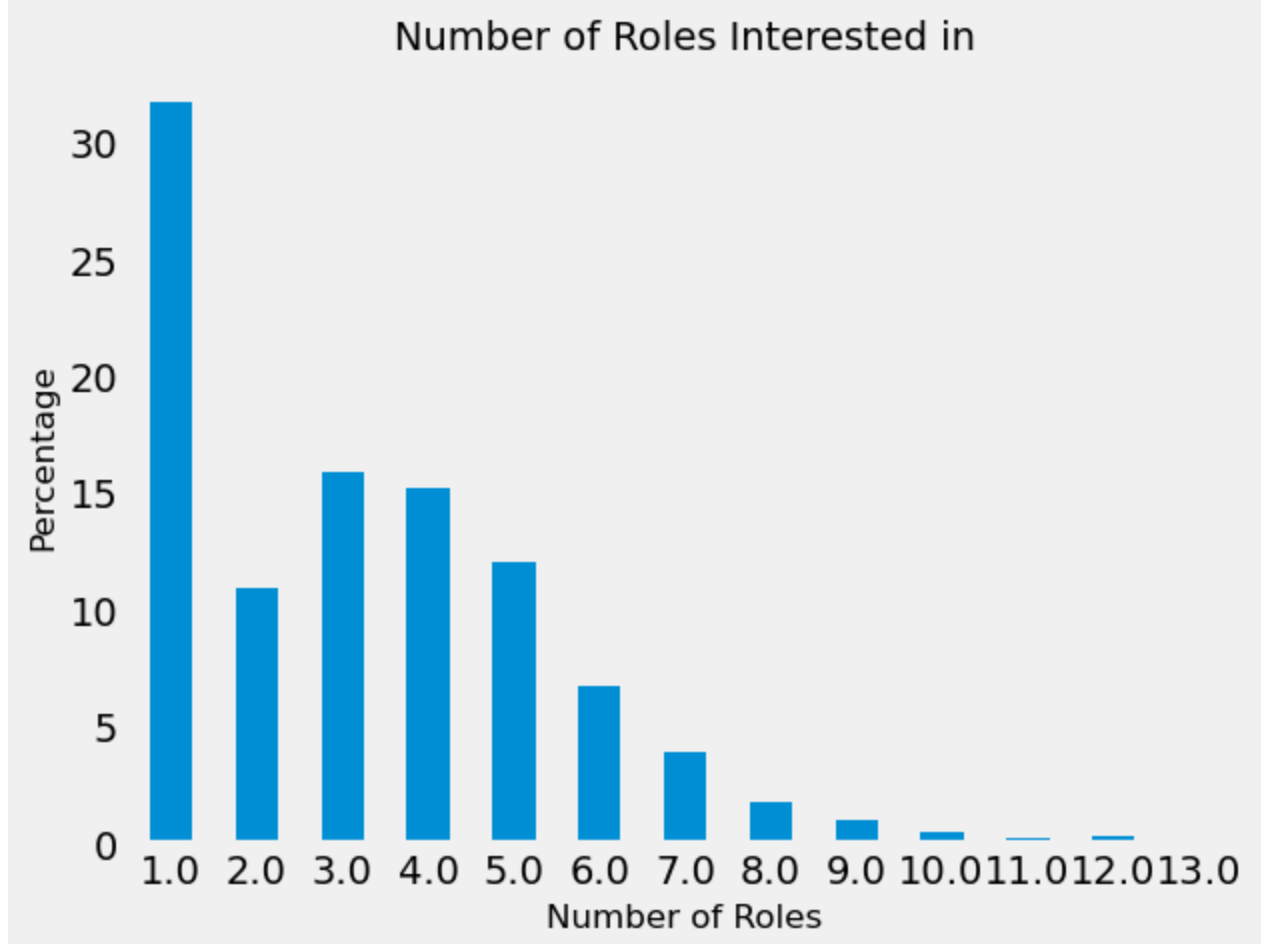
Name: JobRoleInterest, dtype: int64

```
In [11]: numberofroles = (survey_2017['JobRoleInterest'].apply(lambda x: len(x) if x is not np.na
numberofroles = round(numberofroles * 100,1)
print(numberofroles)
```

1.0	31.7
3.0	15.9
4.0	15.2
5.0	12.0
2.0	10.9
6.0	6.7
7.0	3.9
8.0	1.8
9.0	1.0
10.0	0.5
12.0	0.3
11.0	0.2
13.0	0.0

Name: JobRoleInterest, dtype: float64

```
In [12]: numberofroles.sort_index().plot.bar()
plt.xticks(rotation=0)
plt.ylabel('Percentage',fontsize=12)
plt.xlabel('Number of Roles', fontsize=12)
plt.title('Number of Roles Interested in', fontsize=14)
plt.grid(False)
plt.show()
```



Comments:

As we can see by the bar graph above over 30% of respondents are interested in 1 role or specialization. After 1 specialization it drops by over 50%.

Only about 16% are interested in 3 roles/specializations.

Respondents looking at 4 or less roles equal 73% of all respondents.

Lets look at the jobs that the respondents are most interested in.

Job roles of interest:

```
In [13]: job_roles = {}

for roles in survey_2017['JobRoleInterest']:
    if roles is not np.nan:
        for role in roles:
            if role in job_roles:
                job_roles[role] += 1
            else:
                job_roles[role] = 1

job_role = pd.DataFrame(job_roles.items(), columns=['Role_Name', 'Count'])

job_role = job_role.set_index('Role_Name')
print(job_role.sort_values('Count', ascending = False).head(10))
```

Role_Name	Count
fullstack web developer	4198

front end web developer	3533
back end web developer	2772
mobile developer	2305
data scientist	1643
game developer	1628
user experience designer	1469
information security	1326
data engineer	1248
devops / sysadmin	927

```
In [14]: job_role.sort_values('Count', ascending = False).tail(20)
```

```
Out[14]:
```

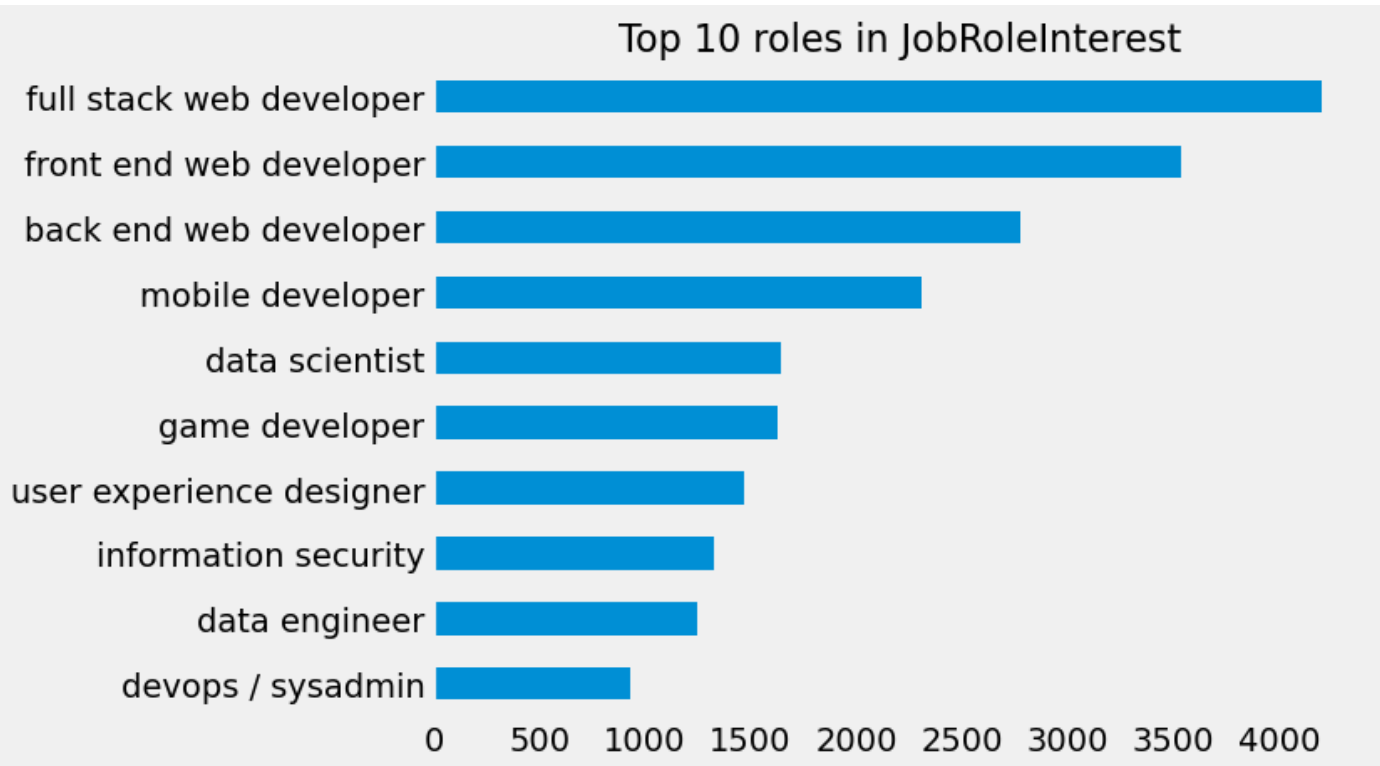
Role_Name	Count
front end web designer	1
computer architect	1
user interface design	1
networking	1
ai and machine learning	1
computer engineer	1
systems programming	1
software engineer (computer science based)	1
technology management	1
ba or developer	1
system engineer	1
technology business liaison	1
network	1
analyst	1
pharmacy tech	1
data journalist / data visualist	1
desings	1
infrastructure architect	1
tech art	1
it specialist	1

Comments:

As we can see by the tail we are losing some answers. A respondent did not necessarily list all roles in the final question as expected or it was not input correctly. We have typos and short hand. This would be alot to sort and filter through to clean up. The good news is that all of the top 10 roles are seperate columns in the survey_2017 dataframe. So I can use this data as well.

```
In [15]: ax = job_role['Count'].sort_values(ascending=False).head(10).plot.barh()
ax.invert_yaxis()
plt.grid(False)
plt.ylabel("")
```

```
plt.title('Top 10 roles in JobRoleInterest', fontsize=16)
plt.show()
```



Comments:

The top 2 courses Legends Code focuses on is web and app development. As we can see above the top 4 roles respondents are interested in is web and mobile development.

Lets count the seperate columns of job interest for the top 4. Combine them into 2 categories web developer and mobile developer and see what the percentage looks like from the total respondents interested.

Job Interest Columns:

```
In [16]: survey_roles = survey_2017.iloc[:,53:66].copy()
```

```
In [17]: Back = survey_roles['JobInterestBackEnd'].sum()
Front = survey_roles['JobInterestFrontEnd'].sum()
Full = survey_roles['JobInterestFullStack'].sum()
mobile = survey_roles['JobInterestMobile'].sum()

Web_Developer = (Back + Front + Full) / 18175
Web_percent = round(Web_Developer * 100,1)

mobile_developer = round((mobile / 18175) * 100,1)

other = round(100 - (mobile_developer + Web_percent),1)

total = Web_percent + mobile_developer

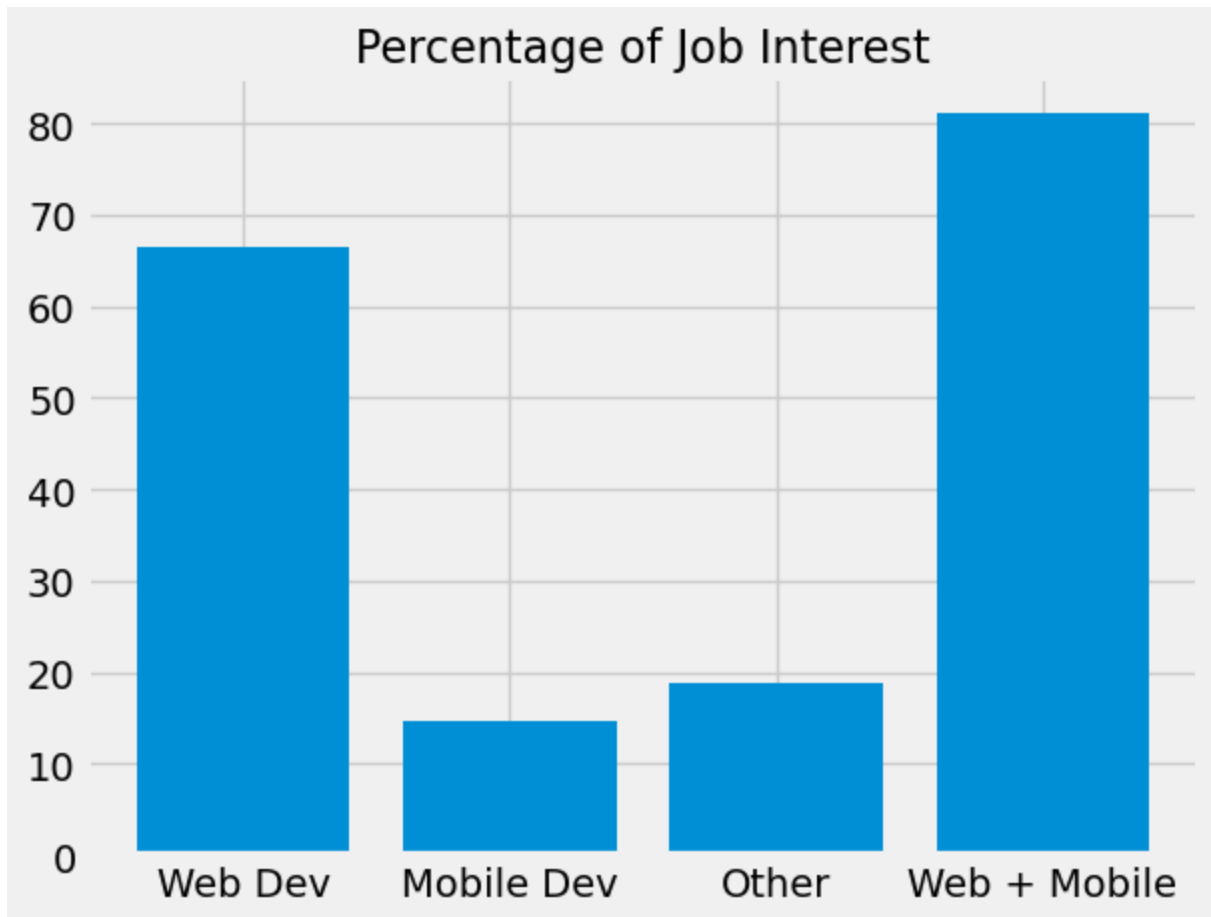
print(Web_percent, mobile_developer, other)

x = ('Web Dev', 'Mobile Dev', 'Other', 'Web + Mobile')
y = (Web_percent, mobile_developer, other, total)

plt.bar(x, y)
plt.ylim(0,85)
```

```
plt.title('Percentage of Job Interest', fontsize=16)
plt.show()
```

66.5 14.7 18.8



Comments:

As we can see above Web Developer is over 60% and mobile developer is over 10%. Combining the two fields gives a total slightly over 80%. This shows that the main two courses are good fields to continue to focus on. With this data right now we dont need to expand our search into other fields. As we can see how low interest is in other fields. The table below is the exact percentages.

Web Developer	Other
66.5%	18.8%
Mobile Developer	
14.7%	
Total	
81.2%	

New Coders Locations and Densities:

Now since we know that people are interested to learn the courses Legends Code offers we need to determine the markets where to advertise. Which market has the most individuals we can reach. We have 2 columns that relate to geographic locations. CountryCitizen and CountryLive. To determine where to advertise I will start with exploring the CountryLive column. The CountryCitizen column is not useful for advertising purposes. We want to reach people where they are currently located.


```
In [18]: survey_2017['CountryLive'].value_counts().head(15)
```

```
Out[18]: United States of America    5791
India                                1400
United Kingdom                       757
Canada                               616
Brazil                               364
Germany                              324
Poland                               265
Russia                               263
Australia                            259
France                               228
Spain                                217
Nigeria                             214
Ukraine                             202
Romania                              171
Italy                                164
Name: CountryLive, dtype: int64
```

```
In [19]: survey_2017['CountryLive'].isnull().value_counts()
```

```
Out[19]: False    15336
True          2839
Name: CountryLive, dtype: int64
```

```
In [20]: #splitting each job interest by country and creating a population count and percentage.
```

```
Back_end = survey_2017['CountryLive'][survey_2017['JobInterestBackEnd'] == 1].value_counts()
Back_end_sum = sum(survey_2017['JobInterestBackEnd'] == 1)
back_end_frequency = round(Back_end / Back_end_sum * 100,2)

Front_end = survey_2017['CountryLive'][survey_2017['JobInterestFrontEnd'] == 1].value_counts()
Front_end_sum = sum(survey_2017['JobInterestFrontEnd'] == 1)
Front_end_frequency = round(Front_end / Front_end_sum * 100,2)

Full_Stack = survey_2017['CountryLive'][survey_2017['JobInterestFullStack'] == 1].value_counts()
Full_Sum = sum(survey_2017['JobInterestFullStack'] == 1)
Full_Stack_Frequency = round(Full_Stack / Full_Sum * 100,2)

Mobile = survey_2017['CountryLive'][survey_2017['JobInterestMobile'] == 1].value_counts()
mobile_sum = sum(survey_2017['JobInterestFrontEnd'] == 1)
mobile_frequency = round(Mobile / mobile_sum * 100,2)
```

```
In [21]: #merging each into its own dataframe
```

```
back_end = pd.merge(Back_end, back_end_frequency, left_index=True, right_index=True)
front_end = pd.merge(Front_end, Front_end_frequency, left_index=True, right_index=True)
full_stack = pd.merge(Full_Stack, Full_Stack_Frequency, left_index=True, right_index=True)
mobile_dev = pd.merge(Mobile, mobile_frequency, left_index=True, right_index=True)

back_end = back_end.rename(columns={'CountryLive_x': 'Population', 'CountryLive_y': 'Percentage'})
front_end = front_end.rename(columns={'CountryLive_x': 'Population', 'CountryLive_y': 'Percentage'})
full_stack = full_stack.rename(columns={'CountryLive_x': 'Population', 'CountryLive_y': 'Percentage'})
mobile_dev = mobile_dev.rename(columns={'CountryLive_x': 'Population', 'CountryLive_y': 'Percentage'})

print('Back End Developer:')
print(back_end)
print()
print('Front End Developer:')
print(front_end)
print()
print('Full Stack Developer:')
print(full_stack)
print()
```

```
print('Mobile Developer')
print(mobile_dev)
```

Back End Developer:

	Population	Percentage
United States of America	1306	40.80
India	198	6.19
United Kingdom	134	4.19
Canada	97	3.03
Germany	56	1.75

Front End Developer:

	Population	Percentage
United States of America	1627	40.20
India	219	5.41
United Kingdom	189	4.67
Canada	145	3.58
Poland	79	1.95

Full Stack Developer:

	Population	Percentage
United States of America	1945	40.26
India	326	6.75
United Kingdom	188	3.89
Canada	137	2.84
Brazil	77	1.59

Mobile Developer

	Population	Percentage
United States of America	1052	25.99
India	173	4.27
United Kingdom	106	2.62
Canada	88	2.17
Brazil	43	1.06

As we can see the United States and India rank the highest in terms of percentage and population. The United Kingdom would be a third option. Now we know for advertising purposes the top 3 Countries are USA, India, UK based on population size.

Spending Money for Learning:

We have an additional column that can help us with our advertising campaign. Above we have the Top 5 countries and we are already eyeing the top 2 or 3 to advertise in. The question now is how much money are people willing to spend on learning?

Here is some information on Legends Code:

- The primary language Legends Code uses is English.
- Legends Code subscription is 59.00 USD a month.

With this information looking at our top 5 countries.

- USA and UK both primary languages is English.
- Canada uses French and English.
- India uses Hindi and English as a secondary.

The 5th country changes in the 4 job interests between:

Brazil who speaks portuguese, Germany speaks German, and Poland speaks polish. For this reason we will stick to the top 4 countries and explore the money poeple are willing to spend on learning.

```
In [22]: #filtering out the top 4 countries needed.

top_4_countries = ['United States of America', 'India', 'United Kingdom', 'Canada']

top_four = survey_2017[survey_2017['CountryLive'].isin(top_4_countries)].copy()

print(top_four['CountryLive'].value_counts())
```

```
United States of America    5791
India                       1400
United Kingdom              757
Canada                      616
Name: CountryLive, dtype: int64
```

```
In [23]: top_four['MoneyForLearning']
```

```
Out[23]: 0          150.0
1           80.0
2        1000.0
5          200.0
6           0.0
...
18156      1000.0
18163         0.0
18164      9500.0
18173         0.0
18174         NaN
Name: MoneyForLearning, Length: 8564, dtype: float64
```

```
In [24]: top_four['MonthsProgramming'].sort_values()
```

```
Out[24]: 1323         0.0
17349         0.0
17350         0.0
17354         0.0
8995          0.0
...
17862         NaN
18052         NaN
18129         NaN
18148         NaN
18163         NaN
Name: MonthsProgramming, Length: 8564, dtype: float64
```

```
In [25]: print(top_four['MonthsProgramming'].isnull().value_counts())
print((top_four['MonthsProgramming'] == 0.0).value_counts())
```

```
False      8247
True        317
Name: MonthsProgramming, dtype: int64
False      8268
True        296
Name: MonthsProgramming, dtype: int64
```

Comments:

I can see already that we have some with money for learning at 0 USD or NaN and months coding with NaN or 0.0.

For the Money for Learning columns we have NaN values. I will change these to 0.0 since we are not sure of the amount of funds available for learning. I don't want to drop the data and rather be more conservative on the available funds for learning.

The people with 0.0 months could be just starting and have not hit one month yet. For these I will change to 1 month. The ones that did not answer I will need to take a closer look of how it will impact the data.

```
In [26]: top_four['MoneyForLearning'].value_counts(dropna = False)
```

```
Out[26]: 0.0      3720
100.0     536
NaN       433
200.0     399
500.0     333
...
28.0       1
765.0      1
103.0      1
538.0      1
4950.0     1
Name: MoneyForLearning, Length: 241, dtype: int64
```

```
In [27]: top_four['MoneyForLearning'] = top_four['MoneyForLearning'].fillna(0.0)

top_four['MoneyForLearning'].value_counts(dropna = False)
```

```
Out[27]: 0.0      4153
100.0     536
200.0     399
500.0     333
50.0      289
...
335.0      1
4100.0     1
32.0       1
6700.0     1
4950.0     1
Name: MoneyForLearning, Length: 240, dtype: int64
```

```
In [28]: nan_months = top_four.loc[top_four['MonthsProgramming'].isnull()]

nan_months['MoneyForLearning'].value_counts().head()
```

```
Out[28]: 0.0      245
500.0      6
200.0      6
100.0      5
2000.0     5
Name: MoneyForLearning, dtype: int64
```

Comments:

As I can see the majority have 0 income that replied with NaN months. For the sake of preserving data and getting the most accurate picture I will keep the NaN months and convert them into 1 as well.

```
In [29]: top_four['MonthsProgramming'] = top_four['MonthsProgramming'].fillna(1)
top_four['MonthsProgramming'] = top_four['MonthsProgramming'].replace([0.0], 1)
```

Median:

```
In [30]: top_four['MonthlySpent'] = top_four['MoneyForLearning'] / top_four['MonthsProgramming']
```

```
top_median = top_four.groupby('CountryLive')['MonthlySpent'].median().sort_values(ascending=True)
display(top_median)
```


```
CountryLive
United States of America    2.777778
Canada                     0.000000
India                      0.000000
United Kingdom              0.000000
Name: MonthlySpent, dtype: float64
```

Mode

```
In [31]: top_four['MonthlySpent'].mode()
```

```
Out[31]: 0    0.0
Name: MonthlySpent, dtype: float64
```

```
In [32]: top_four.groupby('CountryLive')['MonthlySpent'].mean().sort_values().plot.barh(label = 'Average Money Spent Learning',
plt.ylabel('CountryLive')
plt.title('Average Money Spent Learning', fontsize=16, color='purple')
plt.axvline(59, linewidth = 1, linestyle = 'dotted', label = 'Legends Code Subscription',
plt.legend()
plt.grid(False)
```



Comments:

Looking at the average spent. The USA spends the most with Canada second.

If we look at the median. The median is at 0 for 3 of the top 4 countries.

The mode is also 0.

All four countries on average meet the required subscription. India is at the lowest with \$59.23.

Lets break down what the percentage of the population for each country is willing to spend 59 USD or more to learn a month. Then we will need to take the population and percentage to see which is the best to invest

in.

```
In [33]: america = top_four[top_four['CountryLive'] == 'United States of America']
america_spent = round((america['MonthlySpent'] >= 59).value_counts(normalize=True) * 100

uk = top_four[top_four['CountryLive'] == 'United Kingdom']
uk_spent = round((uk['MonthlySpent'] >= 59).value_counts(normalize=True) * 100,2)

canada = top_four[top_four['CountryLive'] == 'Canada']
canada_spent = round((canada['MonthlySpent'] >= 59).value_counts(normalize=True) * 100,2)

india = top_four[top_four['CountryLive'] == 'India']
india_spent = round((india['MonthlySpent'] >= 59).value_counts(normalize=True) * 100,2)

print(america_spent, uk_spent, canada_spent, india_spent)
```

```
False      78.36
True       21.64
Name: MonthlySpent, dtype: float64 False      86.79
True       13.21
Name: MonthlySpent, dtype: float64 False      82.63
True       17.37
Name: MonthlySpent, dtype: float64 False      90.36
True        9.64
Name: MonthlySpent, dtype: float64
```

```
In [47]: fig, ((ax1,ax2),(ax3,ax4)) = plt.subplots(nrows=2, ncols=2, figsize=(7,7))

ax1.pie(america_spent,autopct='%1.1f%%',shadow=True, startangle=90, explode=(0.2,0), col
ax1.set_title('America', fontsize=14)


ax2.pie(canada_spent, autopct='%1.1f%%', shadow=True, startangle=180, explode=(0.2,0), c
ax2.set_title('Canada',fontsize=14)

ax3.pie(uk_spent,autopct='%1.1f%%', shadow=True, startangle=90, explode=(0.2,0), colors=
ax3.set_title('UK',fontsize=14)

ax4.pie(india_spent,autopct='%1.1f%%', shadow=True, startangle=180, explode=(0.2,0), col
ax4.set_title('India',fontsize=14)

plt.text(-4.25,5, 'Percentage of people that spend $59 or more a month', fontsize=14, fo
plt.show()
```

Percentage of people that spend \$59 or more a month



Comments:

As we can see from the pie chart above the percentage of people willing to spend 59 USD will lead us to conclude America and Canada would be our top two picks. 21.6% of respondents in America are willing to spend 59 dollars or above on learning. Canada comes in 2nd place with 17.4% of respondents.

Is this a good representation of data? For example say we wanted to split the advertising budget focusing more resources on the two biggest populations willing to spend money meeting Legends Code monthly subscription price.

Based on percentage we would say America and Canada have the largest portions of students willing to pay 59 USD or more. The question remains what is the size of the population and what do these percents actually add up to?

```
In [35]: uk_spent_numbers = (uk['MonthlySpent'] >= 59).value_counts()  
america_spent_numbers = (america['MonthlySpent'] >= 59).value_counts()  
canada_spent_numbers = (canada['MonthlySpent'] >= 59).value_counts()  
india_spent_numbers = (india['MonthlySpent'] >= 59).value_counts()
```

```
data = { 'True': [1253,107,100,135],
         'False':[4538,509,657,1265]}

raw_numbers = pd.DataFrame(data, index=['United States','Canada', 'U.K.', 'India'])

print(raw_numbers)
```

	True	False
United States	1253	4538
Canada	107	509
U.K.	100	657
India	135	1265

Comments:


As we can see this data is a bit more useful if we wanted to distribute advertising funds unevenly. We are dealing with a small population in each country but India has 28 more individuals than Canada that we can potentially capture.

Dealing with Extreme Outliers

Previously looking through the data of money available for learning I noticed some really high budgets. I'll visualize the data by a box plot for each country and view the monthly budgets. See how extreme the outliers are.

```
In [36]: fig, ((ax1,ax2),(ax3,ax4)) = plt.subplots(nrows=2, ncols=2, figsize=(7,7))

ax1.boxplot(america['MonthlySpent'])
ax1.set_title('USA', fontsize=16)
ax2.boxplot(canada['MonthlySpent'])
ax2.set_title('Canada', fontsize=16)
ax3.boxplot(uk['MonthlySpent'])
ax3.set_title('UK', fontsize=16)
ax4.boxplot(india['MonthlySpent'])
ax4.set_title('India', fontsize=16)
plt.show()
```

Comments:

As we can see above:

- The USA has some learning budgets exceeding 20,000 USD with max of 80,000 USD.
- Canada's learning budgets has a few outliers I would say pretty much 1,500 USD or above with a max of 5,000.
- The UK's learning budget seems to have the same as Canada. Above 1,500 with a max of 8,000.
- India's learning budget seems to be a similar outcome as Canada and the UK. With 1500 USD and above. With a max of 10,000.

Lets take a more closer zoomed in look of each:


```
In [37]: usa_20 = america[america['MonthlySpent'] <= 20000]
canada_20 = canada[canada['MonthlySpent'] <= 2000]
uk_20 = uk[uk['MonthlySpent'] <= 2000]
india_20 = india[india['MonthlySpent'] <= 2000]
```

```
In [38]: fig, ((ax1,ax2),(ax3,ax4)) = plt.subplots(nrows=2, ncols=2, figsize=(7,7))
```

```

ax1.boxplot(usa_20['MonthlySpent'])
ax1.set_title('USA', fontsize=16)
ax2.boxplot(canada_20['MonthlySpent'])
ax2.set_title('Canada', fontsize=16)
ax3.boxplot(uk_20['MonthlySpent'])
ax3.set_title('UK', fontsize=16)
ax4.boxplot(india_20['MonthlySpent'])
ax4.set_title('India', fontsize=16)
plt.show()

```



Comments:

Looking above I would safely say anything above:

- 15,000 USD for the USA is probably an outlier value with smaller percentage
- For the UK and Canada I would say anything above 1,500 USD its probably an outlier and smaller percentage.
- India I would say anything above 1,000 USD is probably an outlier.

Lets see how many outliers we have in each country

```
In [39]: over_15 = america[america['MonthlySpent'] > 15000]
over_15.head(10)
```

Out[39]:

	Age	AttendedBootcamp	BootcampFinish	BootcampLoanYesNo	BootcampName	BootcampRecommend
1393	19.0	0.0	NaN	NaN	NaN	NaN
1964	22.0	0.0	NaN	NaN	NaN	NaN
2478	38.0	1.0	0.0	1.0	NaN	1.0
3168	25.0	1.0	1.0	0.0	App Academy	1.0
3294	28.0	1.0	1.0	0.0	App Academy	1.0
14024	45.0	0.0	NaN	NaN	NaN	NaN
15631	27.0	0.0	NaN	NaN	NaN	NaN
16436	33.0	1.0	1.0	1.0	App Academy	1.0
16650	29.0	0.0	NaN	NaN	NaN	NaN
17093	25.0	0.0	NaN	NaN	NaN	NaN

```
In [40]: over_1500 = canada[canada['MonthlySpent'] > 1500]
over_1500.head(10)
```

Out[40]:

	Age	AttendedBootcamp	BootcampFinish	BootcampLoanYesNo	BootcampName	BootcampRecommend
24	26.0	1.0	1.0	0.0	HackerYou	1.0
6434	25.0	1.0	1.0	0.0	Lighthouse Labs	1.0

6590	20.0	1.0	1.0	0.0	Lighthouse Labs	1.0
7193	22.0	0.0	NaN	NaN	NaN	NaN
8726	39.0	0.0	NaN	NaN	NaN	NaN
13659	24.0	1.0	0.0	0.0	Bloc.io	1.0
14422	24.0	0.0	NaN	NaN	NaN	NaN
16921	45.0	0.0	NaN	NaN	NaN	NaN

```
In [41]: over_15uk = uk[uk['MonthlySpent'] > 1500]
over_15uk.head(10)
```

	Age	AttendedBootcamp	BootcampFinish	BootcampLoanYesNo	BootcampName	BootcampRecommend	(
6483	30.0	0.0	NaN	NaN	NaN	NaN	
7799	25.0	0.0	NaN	NaN	NaN	NaN	
9310	28.0	1.0	1.0	0.0	Makers Academy	1.0	
9401	31.0	0.0	NaN	NaN	NaN	NaN	
16613	49.0	1.0	0.0	0.0	Code Institute	1.0	
16958	34.0	1.0	1.0	1.0	Ace Hacker Academy	1.0	

```
In [42]: over_ind = india[india['MonthlySpent'] > 1500]
over_ind.head(10)
```

Out[42]:

	Age	AttendedBootcamp	BootcampFinish	BootcampLoanYesNo	BootcampName	BootcampRecommend	
	1728	24.0	0.0	NaN	NaN	NaN	NaN
	1755	20.0	0.0	NaN	NaN	NaN	NaN
	2044	21.0	0.0	NaN	NaN	NaN	NaN
	6913	21.0	1.0	0.0	1.0	NaN	1.0
	7989	28.0	0.0	NaN	NaN	NaN	NaN
	8126	22.0	0.0	NaN	NaN	NaN	NaN
	9410	38.0	0.0	NaN	NaN	NaN	NaN
	12451	24.0	0.0	NaN	NaN	NaN	NaN
	13398	19.0	0.0	NaN	NaN	NaN	NaN
	15587	27.0	0.0	NaN	NaN	NaN	NaN

Comments:

In the USA out of the 10 outliers above 15,000 3 are attending a boot camp.

True outliers: 1393,1964,2478,14024,15631,16650,17093

In Canada we have 8 outliers above 1,500 USD. 4 of them are currently attending a boot camp.

True outliers: 7193,8726,,14422,16921

In the U.K. we have 6 outliers above 1,500 USD. 3 are attending a bootcamp.

True outliers: 6483,7799,9401

In India we have 10 outliers and only 1 attended a bootcamp.

True outliers: 1728,1755,2044,7989,8126,9410,12451,13398,15587

Removing outliers and viewing impact:

```
In [43]: america_fresh = america.drop([1393,1964,2478,14024,15631,16650,17093])


canada_fresh = canada.drop([7193,8726,14422,16921])

uk_fresh = uk.drop([6483,7799,9401])

india_fresh = india.drop([1728,1755,2044,7989,8126,9410,12451,13398,15587])
```

```
In [44]: fig, ((ax1,ax2),(ax3,ax4)) = plt.subplots(nrows=2, ncols=2, figsize=(7,7))

ax1.boxplot(america_fresh['MonthlySpent'])
ax1.set_title('USA', fontsize=16)
ax2.boxplot(canada_fresh['MonthlySpent'])
ax2.set_title('Canada', fontsize=16)
ax3.boxplot(uk_fresh['MonthlySpent'])
ax3.set_title('UK', fontsize=16)
ax4.boxplot(india_fresh['MonthlySpent'])
ax4.set_title('India', fontsize=16)
plt.show()
```



```
In [45]: america_fresh['MonthlySpent'].mean().round()
canada_fresh['MonthlySpent'].mean().round()
```

```
uk_fresh['MonthlySpent'].mean().round()  
india_fresh['MonthlySpent'].mean().round()
```

Out[45]: 32.0

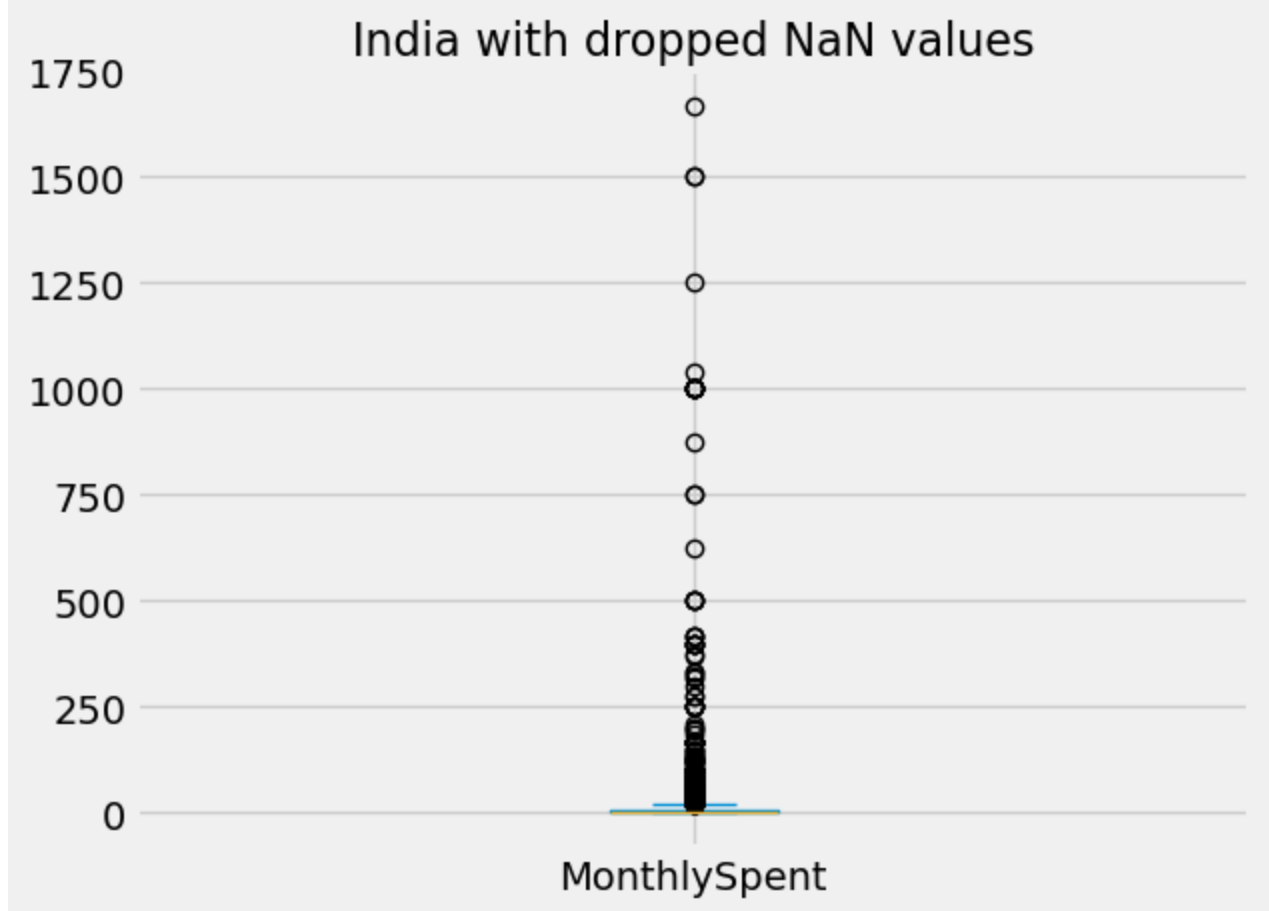
Comments:

Country	Average	Previous Average
USA	172	229
Canada	106	128
U.K.	56	67
India	32	59

We can see how the mean for each country has drastically decreased after removing the top outliers. I took a conservative approach originally. To recall the data in the MoneyforLearning column had many NaN values. What I originally did was converted these NaN values to a 0.00 amount. This altered my analysis throughout the project. As it stands right now the USA and Canada would be my top 2 choices to market in. India was showing potential originally having a larger population than Canada. What I will do now is go back to the original data and drop the NaN values in India. I will then remove the outliers like I did and see how this impacts the average monthly spent column.

Dropping NaN values in India instead of converting to 0.00:

```
In [46]: #capturing only the country India  
  
india_country = ['India']  
  
top_india = survey_2017[survey_2017['CountryLive'].isin(india_country)].copy()  
  
# dropping NaN values  
top_india = top_india.dropna(subset='MoneyForLearning')  
top_india = top_india.dropna(subset='MonthsProgramming')  
  
#converting 0 months to 1 month  
  
top_india['MonthsProgramming'] = top_india['MonthsProgramming'].replace([0.0], 1)  
  
#recreating the monthly spent column  
  
top_india['MonthlySpent'] = top_india['MoneyForLearning'] / top_india['MonthsProgramming']  
  
#confirming same outliers  
  
over_india = india[india['MonthlySpent'] > 1500]  
  
#removing outliers  
  
top_india = top_india.drop([1728,1755,2044,7989,8126,9410,12451,13398,15587])  
  
top_india['MonthlySpent'].plot.box()  
plt.title('India with dropped NaN values', fontsize=16)  
plt.show()  
  
top_india['MonthlySpent'].describe()
```



```
Out[46]: count    1212.000000
mean       34.841987
std        142.638259
min         0.000000
25%         0.000000
50%         0.000000
75%         8.103070
max        1666.666667
Name: MonthlySpent, dtype: float64
```

Comments:

As we can see this barely impacted the results. The average went from 32 USD a month to 34.84 USD a month. The only way to increase the average would be to keep more of the outliers. I only removed the top 9 outliers and this dropped India's average from 59.23 USD to a range of 32-35 USD. Too much weight is placed on the top 9 outliers so I do not think this would be a wise decision to include them.

Recommendation:

My recommendation for Legends Code would be to run an advertising campaign in the United States of America and Canada. The reasoning behind the top 2.

- USA has a monthly average of 172 USD to spend on learning
- Canada has a monthly average of 106 USD to spend on learning
- Both countries are interested in Web and App development.
- Both countries use the English language.

For the disbursement of the advertising budget:

I broke down the interest of each programming field by job role originally. The United States of America by far had the the highest percentages. Around 40% interested in Web development and 25% interested in app

development.

Canada even though was in the top 5 countries had a much lower percentage around 3% for each job type.

As for population the United States had 12 times the population size of Canada.

- Given this information I would primarily focus the advertising campaign in the United States.

If I were to weigh it I would put 80% of the advertising budget in the US and 20% of the advertising budget in Canada. The weight of distribution of advertising can be left for Legends Code executive team to decide.

I would not rule out Canada but I would give a portion of the advertisement budget and see how it generates traffic to Legends Code's website.

My recommendation for a future survey:

We had much more of a population size from the United States of America compared to the other countries. I would recommend setting aside a future budget to survey people in Canada and India directly. Both of these countries still have potential to be a good source of new programmers. With Canada we had a small population but the monthly budget for learning is there. In India we had a decent size population but the monthly budget was low. We were also missing a good portion of a monthly budget for learning in India.